



## ABOUT ME

Hello, I'm Lu, an Assistant Professor at the University of Surrey and a long-term collaborator with the Visual Informatics Group (VITA) at UT Austin. My general research interests focus on understanding and leveraging low-dimensionality in machine learning, impacting areas such as **efficient training and inference of large foundation models, understanding and enhancing reasoning in LLMs, and developing hardware-friendly machine learning algorithms**. I also pursue interdisciplinary applications in geospatial analytics, medical imaging, and healthcare. I am committed to producing transformative research of the highest caliber.

## RESEARCH INTERESTS

# Efficient and Scalable Foundation Models  
(Computational efficiency and inference optimization)

# Understanding and Enhancing NNs  
(performance, robustness, and reasoning capabilities)

# Interdisciplinary AI Applications  
(Integration with Geospatial Analytics, Medical Imaging, and others)

## PROFESSIONAL EXPERIENCE

University of Surrey  
06/2024 – Present

### Assistant Professor

- School of Computer Science and Electronic Engineering  
Nature Inspired Computing and Engineering Group

Eindhoven University of Technology  
07/2023 – 12/2023

### Postdoctoral Researcher

- Department of Mathematics and Computer Science  
Data and AI cluster

Google, New York Office  
07/2023 – 09/2023

### AI Researcher (Intern)

- Build Efficient Large Language Models (LLM)

## EDUCATION

Eindhoven University of Technology  
10/2018 - 2/2023

### Ph.D in Computer Science

**Department:** Mathematics and Computer Science

**Specialization:** Knowledge Elicitation, Data Efficiency, Model Efficiency

**Promoters:** Prof. Dr. Mykola Pechenizkiy; Dr. Vlado Menkovski

Harbin Institute of Technology  
(Shenzhen)  
09/2015 - 07/2018

### Master in Control Engineering

**Department:** Mechanical Engineering and Automation

**Specialization:** Machine Learning, Robotics

**Promoters:** Prof. Dr. Xiaorui Zhu

Harbin Institute of Technology  
09/2009 - 07/2013

### Bachelor in Electrical Engineering and Automation

**Department:** Information and Electrical Engineering

## AWARDS AND HONOURS

---

- 12/2022 Best Paper Award at Learning on Graphs Conference (LoG). 2022.
- 06/2017 Best Paper Nomination Award at International Conference on Computer Vision Systems (ICVS), 2017

## GRANTS

---

### EuroHPC Grants for Computing Time

- Funding Body: European High-Performance Computing Joint Undertaking
- 144000 local core hours on Leonardo BOOSTER (CINECA, Italy).
- Duration: Sep 2025 - Feb 2026
- Role on the Grant: Pi

### NWO Grants for Computing Time

- Funding Body: The Dutch Research Council (NWO)
- Value of Award: 90,431.5 \$
- Duration: April 2023 - April 2025
- Role on the Grant: Co-pi (with Mykola Pechenizkiy and Shiwei Liu)

### NWO Grants for Computing Time

- 2022 EINF-2694: HPC Cloud (CPU): 50.000 hr, HPC Cloud (GPU: NVIDIA GeForce RTX 2080 Ti): 10,000 hr
- 2022 EINF-2943: NVIDIA A100, 1,000,000 Credits (7,812 hr )
- 2023 EINF-5205: HPC Cloud (GPU: NVIDIA GeForce RTX 2080 Ti): 10,000 hr
- 2023 EINF-5206: NVIDIA A100, 1,000,000 Credits (7,812 hr )

## SUPERVISION ACTIVITIES

### Ph.D Projects

- 2024, Robustness of Large Foundation Models - Kappiyath Adarsh, University of Surrey,
- 2024, Resource Efficient 3D World Understanding - Thengane Vishal, University of Surrey,
- 2024, Efficient LLM - Jiayi Li, University of Surrey (Co-supervising with Dr. Xilu Wang)

## TEACHINGS

---

- |  |      |
|--|------|
| • Deep Learning (2AMM10) as TA, Eindhoven University of Technology | 2020 |
| • Deep Learning (2AMM10) as TA, Eindhoven University of Technology | 2021 |
| • Deep Learning (2AMM10) as TA, Eindhoven University of Technology | 2023 |
| • Practical Business Analytics (Com3018), University of Surrey     | 2024 |
| • Deep Learning and Advanced AI (COM3025), University of Surrey    | 2025 |

## SERVICE

---

- 2025 : Conference Reviewer: CPAL, ICLR, CVPR, BMVC, NeurIPS
- 2024: Conference Reviewer: NeurIPS, DAC, UAI, ICML, CPAL  
Journal Reviewer: ACM Computing Surveys
- 2023 : Conference Reviewer: NeurIPS, UAI, ICLR SNN(workshop)  
Journal Reviewer: Clinical Epidemiology
- 2022: Conference Reviewer: SNN
- 2020: Conference Reviewer: ECML-PKDD

## INVITED TALKS

---

- 2024: The power of model sparsity, Multimedia Analytics (MA) Laboratory at City University of Hong Kong
- 2023: LLM pruning, Visual Informatics Group @ University of Texas at Austin,  
Meta universe and Digital Human @ AI time PhD Debate, Tsinghua
- 2022: Model/supervision Efficiency at Xu Lab, Carnegie Mellon University,
- 2020: Going beyond training ML models with labels at EDGE AI, Eindhoven University of Technology,

## ORGANIZATIONAL CONTRIBUTIONS

---

- NeurIPS 2024 Challenge: Edge-Device Large Language Model Competition
- IEEE CAI 2025 Workshop: Stable Training Paradigms for LLMs: Reducing Instability, Increasing Capacity
- IEEE CAI 2025 Workshop: Secure, Private, and Fair Federated Optimization and Learning

## RESEARCH & SELECTED PUBLICATIONS

---

### Overall: (as of June 2025)

- **Over 50** papers (20A\* and 5 A top AI conference papers, CORE Conference Ranking), 1 Journal Paper

### Highlights:

- **1 Best Paper Award, 1 Best Paper Nomination Award.**
- **7 ICML, 4 ICLR, 6 NeurIPS, 2 EMNLP, 2 Interspeech, 1 BMVC, 1 ICASSP, 1 AAAI, 1 UAI. 1 ACL**

### Selected Publications

- **Lu Yin**, You Wu, .etc. *Outlier Weighed Layerwise Sparsity (OWL): A Missing Secret Sauce for Pruning LLMs to High Sparsity*. The Forty-first International Conference on Machine Learning (**ICML**), 2024. [Link](#) (Click me)
- **Lu Yin**, Ajay Jaiswal, .etc. *Pruning Small Pre-Trained Weights Irreversibly and Monotonically Impairs "Difficult" Downstream Tasks in LLMs*. The Forty-first International Conference on Machine Learning (**ICML**), 2024. [Link](#)
- **Lu Yin**, Gen Li, Meng Fang, Li Shen, Tianjin Huang, Zhangyang Wang, Vlado Menkovski, Xiaolong Ma, Mykola Pechenizkiy, Shiwei Liu. *Dynamic Sparse Training Is also A Structure Sparsity Learner*. Conference on Neural Information Processing Systems (**NeurIPS**), 2023. [Link](#)
- **Lu Yin**, Shiwei Liu, Fang Meng, Tianjin Huang, Vlado Menkovski, Mykola Pechenizkiy. *Lottery Pools: Winning More by Interpolating Tickets without Increasing Training or Inference Cost*. Thirty-Seventh AAAI Conference on Artificial Intelligence (**AAAI**), 2023. [Link](#)
- **Lu Yin**, Vlado Menkovski, Meng Fang, Tianjin, Huang, Yulong Pei, Mykola Pechenizkiy, Decebal Constantin Mocanu, Shiwei Liu. *Superposing Many Tickets into One: A Performance Booster for Sparse Neural Network Training*. The 38th Conference on Uncertainty in Artificial Intelligence (**UAI**). 2022. [Link](#)
- Adarsh Kappiyath, Abhra Chaudhuri, AJAY KUMAR JAISWAL, Ziquan Liu, Yunpeng Li, Xiatian Zhu, **Lu Yin\***. *SEBRA : Debiasing through Self-Guided Bias Ranking*. The Thirteenth International Conference on Learning Representations. (**ICLR**) 2025. \*corresponding author. [Link](#)

- Di He, Ajay Jaiswal, Songjun Tu, Li Shen, Ganzhao Yuan, Shiwei Liu, **Lu Yin\***. *AlphaDecay: Module-wise Weight Decay for Heavy-Tailed Balancing in LLMs*. Conference on Neural Information Processing Systems (**NeurIPS**), 2025. *\*corresponding author*. [Link](#)
- Tianhao Chen, Xin Xu, Zijing Liu, Pengxiang Li, Xinyuan Song, Ajay Kumar Jaiswal, Fan Zhang, Jishan Hu, Yang Wang, Hao Chen, Shizhe Diao, Shiwei Liu, Yu Li, **Lu Yin\***, Can Yang. *GPAS: Accelerating Convergence of LLM Pretraining via Gradient-Preserving Activation Scaling*. Conference on Neural Information Processing Systems (**NeurIPS**), 2025. *\*corresponding author*. [Link](#)
- Pengxiang Li\*, **Lu Yin\***, Shiwei Liu. *Mix-LN: Unleashing the Power of Deeper Layers by Combining Pre-LN and Post-LN*. The Thirteenth International Conference on Learning Representations. (**ICLR**) 2025 *\*equal contribution*. [Link](#).
- Pengxiang Li\*, **Lu Yin\***, Shiwei Liu. *Outlier-weighted Layerwise Sampling for LLM Fine-tuning*. The 63rd Annual Meeting of the Association for Computational Linguistics. (**ACL Findings**) 2025 *\*equal contribution*. [Link](#).
- Yuxiang Guo, **Lu Yin**, Bo Jiang, Jiaqi Zhang. *TODO: Enhancing LLM Alignment with Ternary Preferences*. The Thirteenth International Conference on Learning Representations. (**ICLR**) 2025. [Link](#)
- Tianjin Huang, **Lu Yin**, Zhenyu Zhang, Li Shen, Meng Fang, Mykola Pechenizkiy, Zhangyang Wang, Shiwei Liu. *Are Large Kernels Better Teachers than Transformers for ConvNets?* International Conference on Machine Learning (ICML), 2023.. [Link](#)
- Jie Ji, Gen Li, **Lu Yin**, .etc. *Advancing Dynamic Sparse Training by Exploring Optimization Opportunities*. The Forty-first International Conference on Machine Learning (**ICML**), 2024. [Link](#)
- Shiwei Liu, **Lu Yin**, Decebal Constantin Mocanu, and Mykola Pechenizkiy. *Do We Actually Need Dense Over-Parameterization? In-Time Over-Parameterization in Sparse Training*. The Thirty-eighth International Conference on Machine Learning (**ICML**), PMLR, 2021. [Link](#)
- Ajay Jaiswal, **Lu Yin**. etc. *FFN-SkipLLM: A Hidden Gem for Autoregressive Decoding with Adaptive Feed Forward Skipping*, Conference on Empirical Methods in Natural Language Processing (**EMNLP**), 2024. [Link](#)
- Abhinav Bandari, **Lu Yin**. etc. *Is C4 Dataset Enough for Pruning? An Investigation of Calibration Data for LLM Pruning*, Conference on Empirical Methods in Natural Language Processing (**EMNLP**), 2024. [Link](#)
- Gen Li, **Lu Yin**, Jie Ji, Wei Niu, Minghai Qin, Bin Ren, Linke Guo, Shiwei Liu, Xiaolong Ma *NeurRev: Train Better Sparse Neural Network Practically via Neuron Revitalization*. The Twelfth International Conference on Learning Representations. (**ICLR**) 2024. [Link](#)
- Vishal Thengane, Jean Lahoud, Hisham Cholakkal, Rao Muhammad Anwer, **Lu Yin**, Xiatian Zhu, Salman Khan. *CLIMB-3D: Continual Learning for Imbalanced 3D Instance Segmentation*. The British Machine Vision Conference (**BMVC**). 2025. [Link](#)
- Wenfang Sun, Xinyuan Song, Pengxiang Li, **Lu Yin**, Yefeng Zheng, Shiwei Liu. *The Curse of Depth in Large Language Models*. Conference on Neural Information Processing Systems (**NeurIPS**), 2025. [Link](#)
- Ajay Jaiswal, Yifan Wang, **Lu Yin**, Shiwei Liu, Runjin Chen, Jiawei Zhao, Ananth Grama, Yuandong Tian, Zhangyang Wang. *From Low Rank Gradient Subspace Stabilization to Low-Rank Weights: Observations, Theories, and Applications*. Forty-Second International Conference on Machine Learning. **ICML**, PMLR, 2025. [Link](#)
- Qiao Xiao, Boqian Wu, **Lu Yin**, Christopher Neil Gadzinski, Tianjin Huang, Mykola Pechenizkiy, Decebal Constantin Mocanu. *Are Sparse Neural Networks Better Hard Sample Learners?* Conference on British Machine Vision Conference. (**BMVC**), 2024. [Link](#)
- Adriana Fernandez-Lopez, Shiwei Liu, **Lu Yin**, Stavros Petridis, Maja Pantic, *Full-Rank No More: Low-Rank Weight Training for Modern Speech Recognition Models*. IEEE International Conference on Acoustics, Speech, and Signal Processing. (**ICASSP**), 2025. [Link](#)
- Boqian Wu, Qiao Xiao, Shiwei Liu, **Lu Yin**, etc. *E2ENet: Dynamic Sparse Feature Fusion for Accurate and Efficient 3D Medical Image Segmentation*. Conference on Neural Information Processing Systems (**NeurIPS**), 2024. [Link](#)
- Zihang Liu, Tianyu Pang, Oleg Balabanov, Chaoqun Yang, Tianjin Huang, **Lu Yin**, Yaoqing Yang, Shiwei Liu. *LIFT the Veil for the Truth: Principal Weights Emerge after Rank Reduction for Reasoning-Focused Supervised Fine-Tuning*. Forty-Second International Conference on Machine Learning. **ICML**, PMLR, 2025. [Link](#)
- Shiwei Liu, Tianlong Chen, Xiaohan Chen, Zahra Atashgahi, **Lu Yin**, Huanyu Kou, Li Shen, Mykola Pechenizkiy, Zhangyang Wang, and Decebal Constantin Mocanu. *Sparse Training via Boosting Pruning Plasticity with Neuroregeneration*. The Thirty-fifth Conference on Neural Information Processing Systems (**NeurIPS**), 2021. [Link](#)
- Zahra Atashgahi, Xuhao Zhang, Neil Kichler, Shiwei Liu, **Lu Yin**, Mykola Pechenizkiy, Raymond Veldhuis, Decebal Constantin Mocanu. *Supervised Feature Selection with Neuron Evolution in Sparse Neural Networks*. Transactions on Machine Learning Research (**TMLR**).. [Link](#)
- Tianjin Huang, Tianlong Chen, Meng Fang, Vlado Menkovski, Jiaxu Zhao, **Lu Yin**, Yulong Pei, Decebal Constantin Mocanu, Zhangyang Wang, Mykola Pechenizkiy, Shiwei Liu. *You Can Have Better Graph Neural Networks by Not Training Weights at All: Finding Untrained GNNs Tickets*. Learning on Graphs Conference (**LoG**). 2022. (**BEST PAPER AWARD**). [Link](#)
- Xiaorui Zhu, **Lu Yin**, Fucheng Deng. *Wind Disturbance Rejection in Position Control of Unmanned Helicopter by Nonlinear Damping*. International Conference on Computer Vision Systems (**ICVS**). Springer, Cham, 2017: 590-599. (**BEST PAPER NOMINEES AWARD**). [Link](#)

**More in:** <https://scholar.google.com/citations?user=G4Xe1NkAAAAJ>